

Modern Approaches for Heuristic Algorithms in Network Architecture Searches (NAS) and tinyML


Daniel Duclos-Cavalcanti

Advisor:

M.Sc. and M.Eng. Alex Hoffman

Wissenschaftliches Seminar VLSI Entwurfsverfahren

TUM SoSe 2022



Contents

Introduction

- Historical Context and Motivation

Background

- Neural Architecture Search (NAS)
- Search Space
- Search Strategies
- Performance Estimation


State of The Art

- Evolutionary Neural Architecture Search (ENAS)
- TinyML and NAS

Conclusion


Introduction

- Deep Learning
 - Groundbreaking technology
 - Algorithms complexity increase
 - Architecture engineering and domain expertise
- Neural Architecture Search (NAS)



Background – Neural Architecture Search (NAS)

- NAS methods categorized by three dimensions:
 - Search Space
 - Search Strategy
 - Performance Estimation Strategy




NAS – Search Space

- Optimization problem
- Exploration x Exploitation trade-off
- Chain-Structured NN (a) x Multi-Branch NN (b)




(a)



(b)

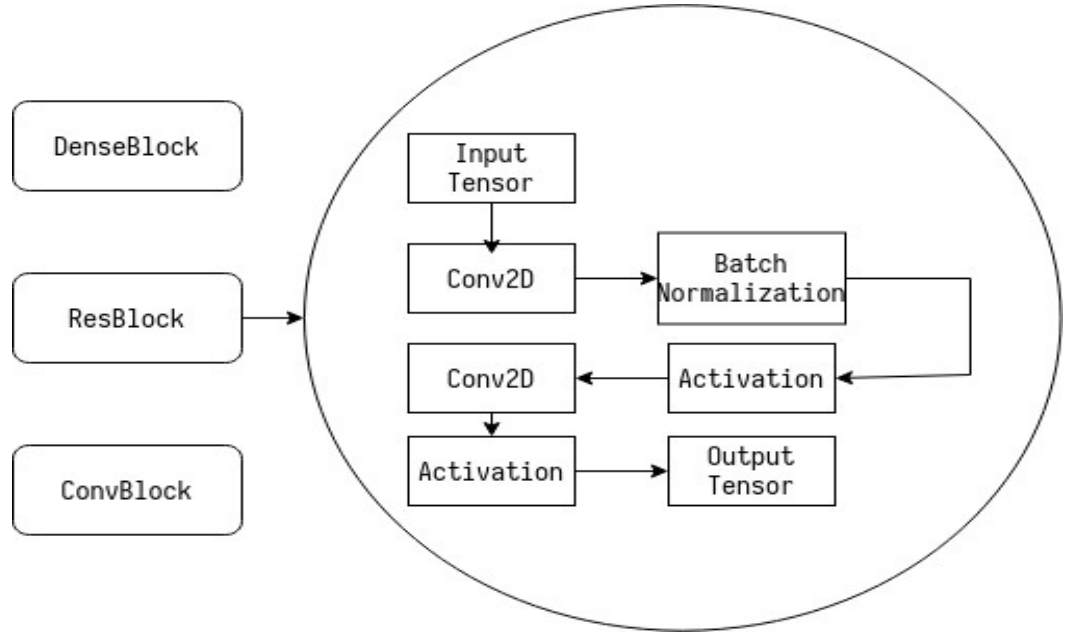
NAS – Search Space

- Optimization problem
- Exploration x Exploitation trade-off
- Cells x Blocks
- Macro x Micro architectures




NAS – Search Space

- Optimization problem
- Exploration x Exploitation trade-off
- Cells x Blocks
- Macro x Micro architectures



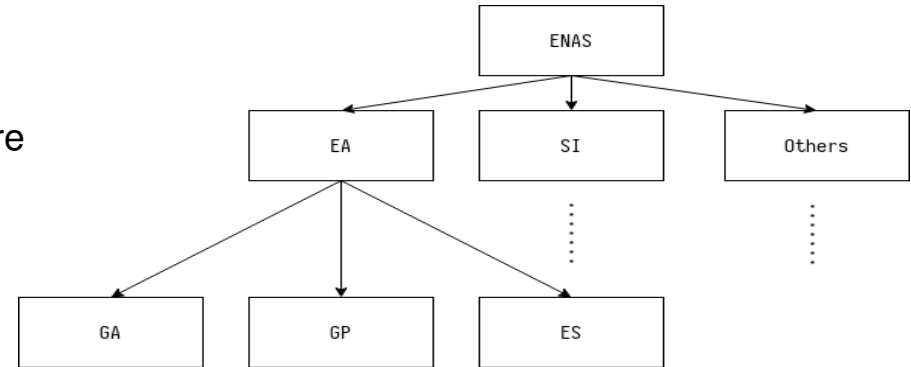
NAS – Search Strategy

- Reinforcement Learning (RL) - Based
 - Agent's action – generation of a neural architecture
 - Action space – search space
 - Agent's reward – performance estimation
- Evolutionary Computation (EC) - Based
- Gradient-Based



NAS – Search Strategy

- Reinforcement Learning (RL) - Based
 - Agent's action – generation of a neural architecture
 - Action space – search space
 - Agent's reward – performance estimation
- Evolutionary Computation (EC) - Based
- Gradient-Based



NAS – Performance Estimation Strategy

- Bias vs Speed trade-off


Speed-up Method	How?
Lower Fidelity Estimates	Shorter training time, training on subset of the data, training on downscaled data or with downscaled models.
Learning Curve Extrapolation	Performance is extrapolated after just a small number of epochs and then decided upon directly. Hyperparameters to predict promising architectures after partial learning. Extrapolate partial learning curves to predict and eliminate sub-optimal architectures.
Weight Inheritance	Weights from previously trained models passed down to new ones.
Weight Sharing or One-Shot Models	All architectures as subgraphs of a supergraph (one-shot model). Weights are shared between specific architectures. Only one-shot is trained.

State-of-the-Art – NAS Overview

- Reinforcement Learning (RL) - Based
 - Costly, hundreds even thousands of GPU days
- Gradient Descent - Based
 - Faster and efficient
 - NAS and gradient-based not completely compatible
- Evolutionary Computation (EC) - Based
 - Easily applicable to non-convex optimization problems
 - Competitive models from trivial initial conditions
 - No human interaction

Evolutionary NAS (ENAS)

- Encoding Space
- Encoding Strategy
- Population Update



ENAS Algorithm Flowchart

ENAS – Encoding Space



- Initial Space + Search Space
- Initialization
 - Trivial Conditions
 - Rich Initialization
 - Random Initialization
- Bias vs Speed

ENAS – Encoding Strategy

- How network architectures are encoded into individuals
- Fixed-length vs variable length
- Easier implementation vs greater discovery freedom

ENAS – Population Update Method

- Selection
 - Elitism
 - Discard worst or oldest
 - Roulette
 - Tournament Selection
- Evolutionary Operators
 - Mutation
 - Crossover



Source: <https://minapecheux.com/website/2018/09/15/a-peek-at-genetic-algorithms/>

State-of-the-Art – tinyML and NAS

- TinyML
 - Port AI algorithms to constrained devices
 - Training, pruning, quantizing and hardware design
 - Not optimal, mutual influence of HW and algorithm
- McuNET
 - TinyEngine and TinyNAS
 - Optimizes search space according to memory, energy consumption, storage and latency
- MicroNETs
 - Tailored differentiable NAS based on DARTS
 - Regularization term, balance between eFlash memory x SRAM x accuracy

Conclusions

- NAS is a complex optimization problem with great promise.
- NAS can be used to optimize architecture search for multi-objective goals.
- This proves useful within the tinyML paradigm:
 - Energy consumption
 - Latency
 - Memory Usage
 - Storage

Questions